


Peter Müller

OWNERSHIP IN PROGRAM VERIFICATION – FROM SEPARATION LOGIC TO RUST AND BACK

ETH zürich



Outline

- The Viper intermediate verification language
 - Reasoning about the heap
 - Abstraction
 - Advanced separation logic in Viper
- Verification of Rust programs
 - Ownership
 - References
- Conclusion



- viper.ethz.ch
- Try online: <http://viper.ethz.ch/tutorial>
- Install as VS Code extension

Heap model: an object-based language

```
field val: Int

method foo() returns (res: Int)
{
  var cell: Ref
  cell := new(val)
  cell.val := 5
  res := cell.val
}
```

- A heap is a set of objects
- No classes: each object has all fields declared in the entire program
 - Type rules of a source language can be encoded
 - Memory consumption is not a concern since programs are not executed
- Objects are accessed via references
 - Field read and update operations
- No explicit de-allocation (garbage collector)
 - Conceptually, objects could remain allocated

The frame problem

```
field f: Int  
field g: Int
```

```
method set(p: Ref, v: Int)  
  requires p != null  
  ensures p.f == v  
{  
  p.f := v  
}
```

```
x.f := 0  
x.g := 0  
set(x, 5)  
assert x.g == 0
```



- Bad idea: inspect body of callee to determine which field locations are modified
 - Not modular
 - Does not work for abstract methods
- Bad idea: assume conservatively that all field locations may be modified
 - Callee needs a specification for all field locations, even those it does not change
 - Not modular: procedure specifications need to change when a new field is declared

Aliasing

```
field val: Int

method foo(p: Ref, q: Ref)
  requires p != null && q != null
  ensures p.val == 5
{
  p.val := 5
  q.val := 7
}
```




- Bad idea: require programmer to explicitly specify all non-aliasing properties
 - Does not work for unbounded data structures
 - Violates information hiding

Further challenges of heap verification

- Writing specifications that preserve [information hiding](#)
- [Concurrency](#), e.g., data races

Access permissions

- Associate each heap location with a permission
- Permissions are created when the heap location is allocated: `new(f, g)`
- Permissions are held by method executions (or loop iterations)
- Read or write access to a memory location requires permission
- Permissions can be transferred, but not duplicated or forged



`x.f := 5`



`y.f := 5`



`z.g := x.f`



`x.f := y.f`



Permission assertions

- Permissions are denoted in assertions by **access predicates**
 - Access predicates are not permitted under negations, disjunctions, and on the left of implications
- Assertions may contain both permissions and value constraints
- Many assertions that occur in a program must be **self-framing**, that is, include all permissions to evaluate the heap accesses in the assertion

acc(p.f)

$p.f \mapsto _$

acc(p.f) && p.f > 0

$\exists V \cdot p.f \mapsto V \wedge V > 0$

requires p.f > 0



Separating conjunction

- To handle aliasing, we introduce a new connective: **separating conjunction**
- $\mathbf{A} * \mathbf{B}$ holds in a state if:
 - both \mathbf{A} and \mathbf{B} hold, and
 - the **sum of the permissions** in \mathbf{A} and \mathbf{B} are held in that state
 - $\mathbf{A} * \mathbf{B}$ and $\mathbf{A} \wedge \mathbf{B}$ are equivalent if \mathbf{A} and \mathbf{B} do not contain access predicates
- Holding permission to locations $p.f$ and $q.f$ implies that p and q do not alias
- Viper's $\&\&$ is separating conjunction
- For the call `swap(x, x)`, the precondition is equivalent to false

$$\text{acc}(p.f) * \text{acc}(q.f) \Rightarrow p \neq q$$

```
method swap(a: Ref, b: Ref)
  requires acc(a.f) && acc(b.f)
```

Framing

Frame rule

$$\frac{\{ \mathbf{A} \} \text{ S } \{ \mathbf{B} \}}{\{ \mathbf{A} * \mathbf{C} \} \text{ S } \{ \mathbf{B} * \mathbf{C} \}}$$

where S does not assign to a local variable that is free in **C**

- The frame **C** must be self-framing
 - If heap locations constrained by **C** are disjoint from those modified by S then **C** is preserved
 - Otherwise, the precondition of the conclusion is equivalent to false (the triple holds trivially)
- Example

$$\frac{\{ \text{acc}(x.f) \} \text{ } x.f := 5 \text{ } \{ \text{acc}(x.f) * x.f = 5 \}}{\{ \text{acc}(x.f) * \text{acc}(y.f) * y.f = 7 \} \text{ } x.f := 5 \text{ } \{ \text{acc}(x.f) * x.f = 5 * \text{acc}(y.f) * y.f = 7 \}}$$

Framing for method calls

```
method set(p: Ref, v: Int)
  requires acc(p.f)
  ensures  acc(p.f) && p.f == v
{
  p.f := v
}
```

```
// assume we have acc(x.f) && acc(y.f)
assume y.f == 7
set(x, 5)
assert x.f == 5 && y.f == 7
```


$$\frac{\frac{\{ \text{acc}(p.f) \} \text{ method set}(p, v) \{ \text{acc}(p.f) * p.f = v \}}{\{ \text{acc}(x.f) \} \text{ set}(x, 5) \{ \text{acc}(x.f) * x.f = 5 \}}}{\{ \text{acc}(x.f) * \text{acc}(y.f) * y.f = 7 \} \text{ set}(x, 5) \{ \text{acc}(x.f) * x.f = 5 * \text{acc}(y.f) * y.f = 7 \}}}$$

- A method modifies at most the heap locations to which it has permission


Permission transfer

```
method set(p: Ref, v: Int)
  requires acc(p.f)
  ensures  acc(p.f) && p.f == v
```

```
{
```




```
  p.f := v
```



```
}
```

```
// assume we have acc(x.f) && acc(y.f)
assume x.f == 2 && y.f == 7
```




```
set(x, 5)
```

Framing!

```
assert x.f == 5 && y.f == 7
```

Permission transfer for method calls



- Permissions are held by **method executions** or loop iterations
- Calling a method **transfers permissions from the caller to the callee** (according to the method precondition)
- Returning from a method **transfers permissions from the callee to the caller** (according to the method postcondition)
- **Residual permissions are framed around the call**

Framing for loops

```
// assume we have acc(x.f) && acc(y.f)
x.f := 0
y.f := 7
while (x.f < 10)
  invariant acc(x.f)
{
  x.f := x.f + 1
}
assert y.f == 7
```

$$\frac{\frac{\{ \text{acc}(x.f) \} \text{ while}(x.f < 10) \{ \dots \} \{ \text{acc}(x.f) * \neg x.f < 10 \}}{\{ \text{acc}(x.f) * \text{acc}(y.f) * y.f = 7 \} \text{ while}(x.f < 10) \{ \dots \} \{ \text{acc}(x.f) * \neg x.f < 10 * \text{acc}(y.f) * y.f = 7 \}}}{\{ \text{acc}(x.f) \} \text{ while}(x.f < 10) \{ \dots \} \{ \text{acc}(x.f) * \neg x.f < 10 \}}$$


Permission transfer: inhale and exhale operations

- **inhale** **A** means:

- obtain all permissions required by assertion **A**
- assume all logical constraints

```
inhale acc(x.f) && x.f == 2
```


2



- **exhale** **A** means:

- assert all logical constraints
- check and remove all permissions required by assertion **A**
- havoc any locations to which all permission is lost

```
exhale acc(x.f) && x.f == 2
```



Simplified proof obligations for method bodies and calls

```
method foo() returns (...)  
  requires A  
  ensures B  
  { S }
```

```
x := foo()
```

▪ Encoding *without heap*

- Body

```
assume A  
// encoding of S  
assert B
```

- Call

```
assert A[...]  
havoc x  
assume B[...]
```

▪ Encoding *with heap*

- Body

```
inhale A  
// encoding of S  
exhale B
```

- Call

```
exhale A[...]  
havoc x  
inhale B[...]
```

▪ **inhale** and **exhale** are permission-aware analogues of **assume** and **assert**

Verifying memory safety

- Memory safety is the absence of errors related to memory accesses, such as, null-pointer dereferencing, access to un-allocated memory, dangling pointers, out-of-bounds accesses, double free, etc.
- Using permissions, Viper verifies memory safety by default

```
var x: Ref  
x.f := 5
```



```
var x: Ref  
x := null  
x.f := 5
```



```
method free(p: Ref)  
  requires acc(p.f)
```

model de-allocation
via method call

```
free(x)  
x.f := 5
```




```
free(x)  
free(x)
```



Challenges revisited

Heap data structures pose three major challenges for sequential verification

- Reasoning about aliasing
 - Permissions and separating conjunction
- Framing, especially for dynamic data structures
 - Sound frame rule, but no support yet for unbounded data structures
- Writing specifications that preserve information hiding
 - Not solved yet



And additional challenges for concurrent programs, e.g., data races

- Permissions are an excellent basis

Outline

- The Viper intermediate verification language
 - Reasoning about the heap
 - Abstraction
 - Advanced separation logic in Viper
- Verification of Rust programs
 - Ownership
 - References
- Conclusion

Limitations of technique as introduced so far

```
field elem: Int
field next: Ref

method head(this: Ref) returns (res: Int)
  requires acc(this.elem)
  ensures  acc(this.elem)
  ensures  res == this.elem
{
  res := this.elem
}
```

- Specification reveals implementation details

```
method append(this: Ref, e: Int)
  requires // permission to all nodes
  ensures  // list was extended
{
  if(this.next == null) {
    var n: Ref
    n := new(elem, next)
    n.next := null
    this.elem := e
    this.next := n
  } else {
    append(this.next, e)
  }
}
```

- Neither permissions nor behavior can be expressed

Predicates

- User-defined predicates consist of a predicate name, a list of parameters, and a self-framing assertion

```
predicate node(this: Ref) {  
  acc(this.elem) && acc(this.next)  
}
```

- Recursive predicates may denote a statically-unbounded number of permissions

```
predicate list(this: Ref) {  
  acc(this.elem) && acc(this.next) &&  
  (this.next != null ==> list(this.next))  
}
```

Static verification with recursive predicates

- A program verifier in general cannot know statically how far to unfold recursive definitions

```
predicate list(this: Ref) {  
    acc(this.next) &&  
    (this.next != null ==> list(this.next))  
}
```

```
inhale list(x)  
y.next := null // do we have permission?
```

Iso-recursive predicates

- An iso-recursive semantics distinguishes between a predicate instance and its body

```
predicate list(this: Ref) {  
  acc(this.elem) && acc(this.next) &&  
  (this.next != null ==> list(this.next))  
}
```

```
inhale list(x)  
x.next := null // no permission
```



- Intuition: permissions are held by method executions, loop iterations, or predicate instances

Folding and unfolding predicates

- Exchanging a predicate instance for its body, and vice versa, is done via fold and unfold statements in the program
- An unfold statement exchanges a predicate instance for its body
- A fold statement exchanges a predicate body for a predicate instance

```
inhale list(x)  
unfold list(x)  
x.next := null
```

```
inhale list(x)  
unfold list(x)  
x.next := null  
fold list(x)  
exhale list(x)
```

- Unfolding-expressions allow one to temporarily unfold a predicate during the evaluation of an expression

Specifying functional behavior

- Using old-expressions and unfolding-expressions, we can specify some aspects of functional behavior
- But: Approach does not work when behavior depends on an unbounded number of fields (e.g., sorting a list)
- And: specifications reveal implementation details

```
method head(this: Ref) returns (res: Int)
  requires list(this)
  ensures list(this)
  ensures res ==
    old(unfolding list(this) in this.elem)
```

Challenges revisited


Heap data structures pose three major challenges for sequential verification

- Reasoning about aliasing
 - Permissions and separating conjunction
- Framing, especially for dynamic data structures
 - Sound frame rule, predicates
- Writing specifications that preserve information hiding
 - Not solved



Data abstraction

- To write implementation-independent specifications, we map the concrete data structure to mathematical concepts and specify the behavior in terms of those



Data abstraction via abstraction functions

- Viper provides **heap-dependent functions**

- side-effect free
- terminating
- deterministic

- Function bodies are **expressions**

- Functions may be **recursive**

- Functions must have a **precondition that frames the function body**, that is, provides all permissions to evaluate the body

```
function content(this: Ref): Seq[Int]
  requires list(this)
{
  unfolding list(this) in
    (this.next == null ?
      Seq[Int]() :
      Seq(this.elem) ++ content(this.next)
    )
}
```

Challenges revisited

Heap data structures pose three major challenges for sequential verification

- Reasoning about aliasing
 - Permissions and separating conjunction
- Framing, especially for dynamic data structures
 - Sound frame rule, predicates
- Writing specifications that preserve information hiding
 - Data abstraction, heap-dependent functions



Outline



- The Viper intermediate verification language
 - Reasoning about the heap
 - Abstraction
 - Advanced separation logic in Viper
- Verification of Rust programs
 - Ownership
 - References
- Conclusion

Limitations of technique as introduced so far

```
field val: Int

method clone(this: Ref) returns (res: Ref)
  requires acc(this.val) // read only
  ensures  acc(this.val) && acc(res.val)
  ensures  this.val == old(this.val)
{
  res := new(val)
  res.val := this.val
}
```

```
predicate person(this: Ref) {
  acc(this.savings) &&
  account(this.savings)
}
```



- To enable framing, methods that only read a data structure must specify that each abstraction remains **unchanged**

- **Sharing** in data structures cannot be expressed (all structures are trees)

Fractional permissions

- To **distinguish read and write access**, permissions can be split and re-combined

- The **permission amount** π is a rational number in $[0;1]$
- Viper syntax allows fractions n/d , **write** for 1, **none** for 0, and **wildcard** for existentially-quantified positive amounts
- **acc**($E.f$) is a shorthand for **acc**($E.f$, **write**), and $P(E)$ for **acc**($P(E)$, **write**)

```
acc(p.f,  $\pi$ )
```

```
acc(P(e),  $\pi$ )
```

- **Field read** requires some **non-zero** permission

```
inhale acc(x.f, 1/2)  
v := x.f
```



- **Field write** requires **full** (write) permission

```
inhale acc(x.f, 1/2)  
x.f := v
```



Manipulating fractional permissions

- Separating conjunction sums up permissions of the conjuncts

`acc(x.f, 1/2) && acc(x.f, 1/2)` is equivalent to `acc(x.f, write)`

- Inhale adds permissions
- Exhale subtracts permissions and *havocs only when **all** permission to a location or predicate instance is removed*
- Values are *framed as long as **some** permission is held*

```
method clone(this: Ref) returns (res: Ref)
  requires acc(this.val, 1/2) // read only
  ensures  acc(this.val, 1/2) && acc(res.val)
  { ... }
```

```
method demo(this: Ref) returns (c: Ref)
  requires acc(this.val)
  {
    var tmp: Int
    tmp := this.val
    c := clone(this) // no havoc of this.val
    assert tmp == this.val
  }
```

Predicates and fractional permissions

- Predicates may contain fractional permissions, for instance, to permit sharing
- **Unfold and fold multiply** the fraction of the predicate with the fractions in the predicate body

```
predicate readCell(this: Ref) {  
    acc(this.cell) && acc(this.cell.val, 1/2)  
}
```


```
inhale acc(readCell(x), 1/4)  
unfold acc(readCell(x), 1/4)  
exhale acc(x.cell.val, 1/8)
```



Sharing in data structures


- Full permissions can describe tree-shaped data structures only

```
predicate person(this: Ref) {  
  acc(this.savings) &&  
  account(this.savings)    }  
}
```




- Fractional permissions allow sharing

```
predicate person(this: Ref) {  
  acc(this.savings) &&  
  acc(account(this.savings), 1/2) }  
}
```



- including unbounded (immutable) sharing

```
predicate person(this: Ref) {  
  acc(this.savings &&  
  acc(account(this.savings), wildcard) }  
}
```



Partial data structures


```
field savings: Ref
predicate account(this: Ref)
predicate person(this: Ref) {
  acc(this.savings) &&
  account(this.savings)
}
```

```
method get(p: Ref) returns (s: Ref)
  requires person(p)
  ensures account(s)
{
  unfold person(p)
  s := p.savings
}
```

```
s := get(p)
s.deposit(100)
p.move() // requires person(p)
```



- To allow clients of `get` to use the person object later, `get` needs to return permissions to the remainder of the object



- Bad idea: return permissions to fields explicitly
 - Violates information hiding
 - Does not work for unbounded structures
- Bad idea: define a dedicated predicate
 - Requires a way to identify the hole
 - Requires ghost code to plug the hole

Separating implication: magic wands

- A magic wand $P \multimap Q$ represents the difference between Q and P



- This allows us to specify our get method

```
method get(p: Ref) returns (s: Ref)
  requires person(p)
  ensures  account(s) && (account(s) --* person(p))
```

- Intuition: permissions are held by method executions, loop iterations, predicate instances, or magic wands

Reasoning with magic wands

- Applying a magic wand

- Viper has a designated statement to apply modus ponens for magic wands


apply $P \multimap Q$

$P * (P \multimap Q) \Rightarrow Q$

- Creating a magic wand

- Viper needs to determine which permissions from the current state need to be moved into the wand such that the wand, together with P , yields Q

package $P \multimap Q$

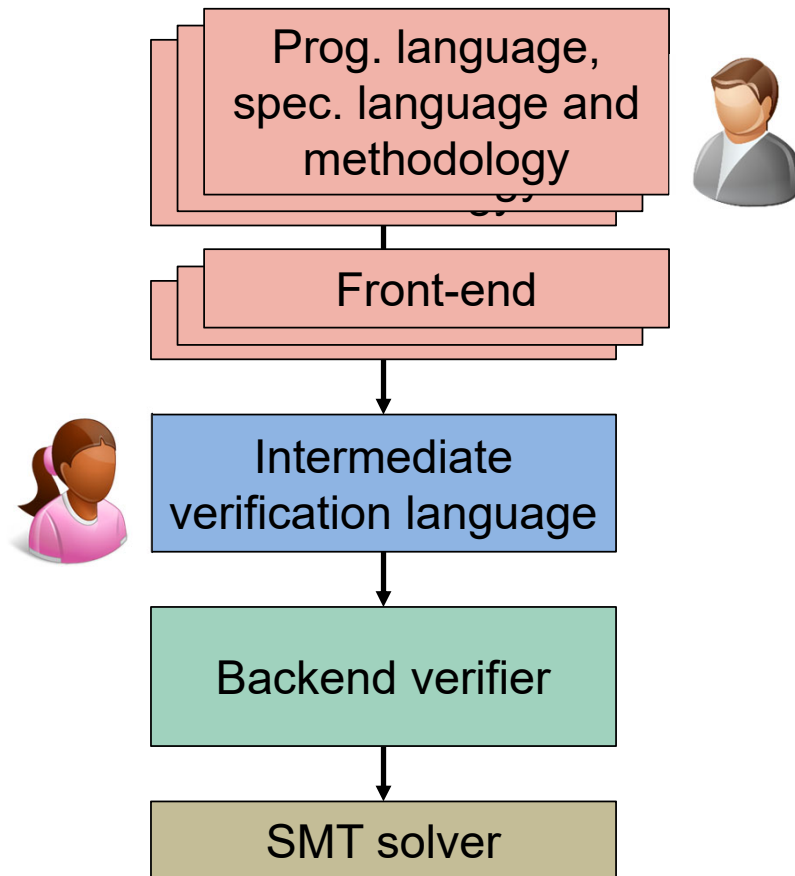


Example revisited

```
method get(p: Ref) returns (s: Ref)
  requires person(p)
  ensures account(s) && (account(s) --* person(p))
{
  unfold person(p)
  s := p.savings
  package account(s) --* person(p) {
    fold person(p)
  }
}
```

```
s := get(p)
s.deposit(100)
apply account(s) --* person(p)
p.move() // requires person(p)
```





gobra

VIPER

Go verification in Gobra

```
requires isList(a) && isList(b)
func client(a *list, b *list) {
    tmp := b.length()
    a.appendBack(100)
    assert b.length() == tmp
}
```

The logo for Gobra, featuring the word "gobra" in a stylized orange font with a small orange icon above the 'o'.

Functional properties

Memory errors

Aliasing

Data races

- To handle these challenges, Gobra uses a verification technique similar to Viper's
- Programmers get exposed to the verification logic
- The overhead is substantial (both amount and complexity of annotations)

Rust and its type system

```
fn client(a: &mut List, b: &mut List)
{
    let tmp = b.length();
    a.appendBack(100);
    assert!(b.length() == tmp);
}
```




Functional properties

No memory errors

Controlled aliasing

No data races

Can we leverage this type system to simplify verification?



$P * \text{rust} \rightarrow * i$

VIPER

Outline

- The Viper intermediate verification language
 - Reasoning about the heap
 - Abstraction
 - Advanced separation logic in Viper
- Verification of Rust programs
 - Ownership
 - References
- Conclusion

Ownership in Rust

```
struct Pair {  
    first: i32,  
    second: i32,  
}  
  
fn demo(mut p: Pair, mut q: Pair) {  
    p.first = 5;  
    q.first = 7;  
    assert!(p.first == 5)  
}
```




- Rust supports ownership via its type system
- An expression representing a memory location (a place) of a type T owns the value in that location
- Owned values cannot be aliased
- Rust leverages this guarantee:
 - To ensure data race freedom
 - To perform automatic memory management (without a garbage collector)


Move assignments

- When reading from a place, **ownership** of the read value **is moved** to the receiving place
 - Except for types that implement the Copy trait, such as `i32`
- When a value is moved out of a place, the place becomes **unusable** until a new value is assigned

```
fn foo(mut p: Pair) {  
    let mut q = p;  
    q.first = 5;  
}
```



```
fn bar(mut p: Pair) {  
    foo(p);  
    p.second = 7;  
}
```



From Rust ownership to permissions

- The notions of ownership in Rust and separation logic are extremely similar

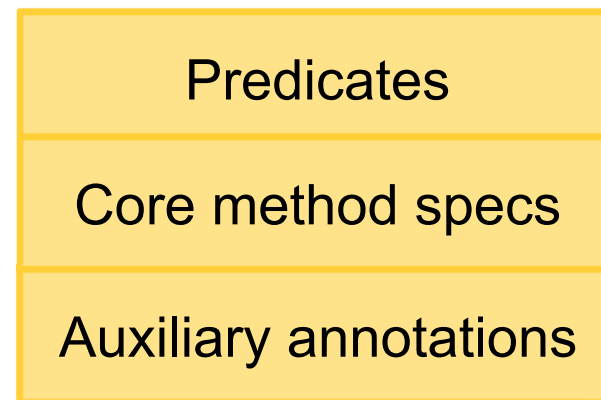

“Rust ownership \approx write permission”

Exclusive ownership

No access through aliases

Ownership transfer upon calls and returns

- Approach: generate a memory safety proof in Viper automatically from Rust types



VIPER

Viper encoding: primitive types

- We model values of primitive types as objects with one field (so that we can later take a reference to the memory location)

```
field val: Int
predicate i32(this: Ref) {
  acc(this.val)
}
```

- Auxiliary operations yield the integer value and create the object for an integer

```
function getInt(this: Ref): Int
  requires i32(this)
{
  unfolding i32(this) in this.val
}
```

```
method makeInt(x: Int) returns (res: Ref)
  ensures i32(res)
  ensures getInt(res) == x
{
  res := new(val);
  res.val := x
  fold i32(res)
}
```

Viper encoding: user-defined types

- The predicate for a struct includes the predicate instances for its fields because ownership in Rust is deep


```
struct Pair {  
  first: i32,  
  second: i32,  
}
```



```
field first: Ref  
field second: Ref  
  
predicate Pair(this: Ref) {  
  acc(this.first) && i32(this.first) &&  
  acc(this.second) && i32(this.second)  
}
```

Viper encoding: methods

```
fn demo(mut p: Pair, mut q: Pair) {  
  p.first = 5;  
  q.first = 7;  
  assert!(p.first == 5)  
}
```



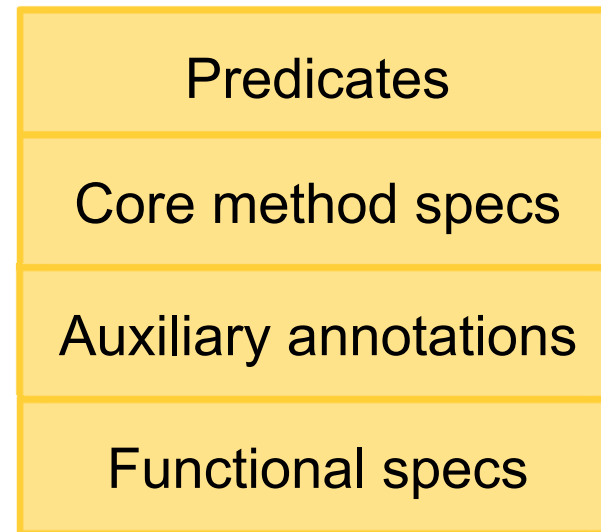

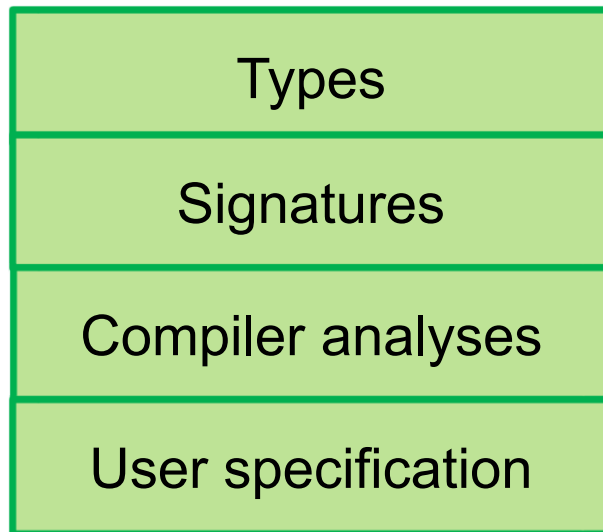
```
method demo(p: Ref, q: Ref)  
  requires Pair(p) && Pair(q)  
{  
  unfold Pair(p)  
  p.first := makeInt(5)  
  unfold Pair(q)  
  q.first := makeInt(7)  
  fold Pair(q)  
  assert getInt(p.first) == 5  
  fold Pair(p)  
}
```



- The encoding generates all annotations necessary to verify the method in Viper

Functional correctness

- The [automatically-generated core proof](#) re-proves memory safety
- The contained permission specifications enable reasoning about [aliasing](#) and [framing](#) (and race freedom in concurrent code)
- Functional specifications can be layered on top easily



Viper encoding: functions and contracts

```
#[pure]
fn max(a: i32, b: i32) -> i32 {
  if a > b { a } else { b }
}
```



```
function max(a: Ref, b: Ref): Ref
  requires i32(a) && i32(b)
  { getInt(a) > getInt(b) ? a : b }
```

```
#[ensures(result.first == old(max(p.first, q.first)))]
fn pmax(p: Pair, q: Pair) -> Pair
{ ... }
```



```
method pmax(p: Ref, q: Ref) returns (res: Ref)
  requires Pair(p) && Pair(q)
  ensures Pair(res)
  ensures unfolding Pair(res) in
    getInt(res.first) ==
      old(unfolding Pair(p) in
        unfolding Pair(q) in
          getInt(max(p.first, q.first)))
  { ... }
```

Outline

- The Viper intermediate verification language
 - Reasoning about the heap
 - Abstraction
 - Advanced separation logic in Viper
- Verification of Rust programs
 - Ownership
 - References
- Conclusion

Borrowing

- With move assignments, functions need to return their arguments in order to transfer ownership back to the caller

```
fn swap(mut p: Pair) -> Pair {  
    let tmp = p.first;  
    p.first = p.second;  
    p.second = tmp;  
    p  
}
```

```
let mut p = Pair{ first: 5, second: 7 };  
p = swap(p);  
println!("{}", p.first);
```

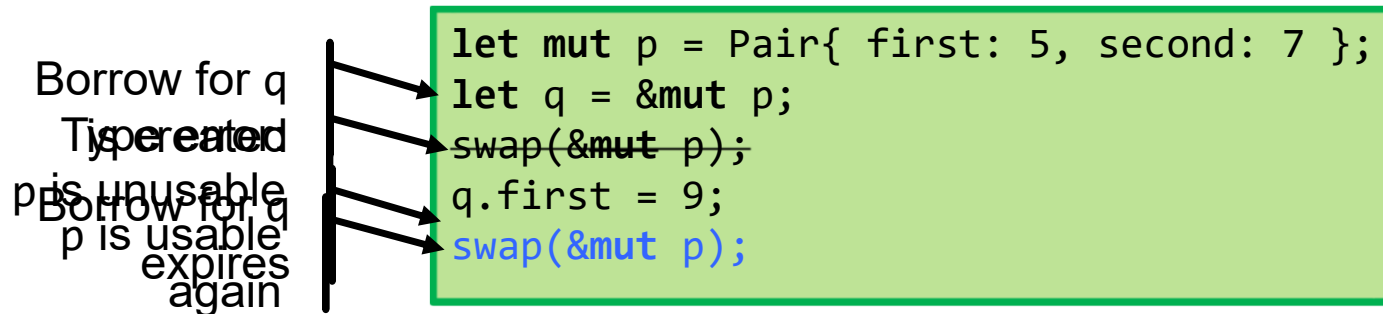
- To avoid transferring ownership back and forth, Rust supports [borrowing](#), which provides [temporary access](#) to a value [through a reference](#)

```
fn swap(p: &mut Pair) {  
    let tmp = p.first;  
    p.first = p.second;  
    p.second = tmp;  
}
```

```
let mut p = Pair{ first: 5, second: 7 };  
swap(&mut p);  
println!("{}", p.first);
```

Borrow checking


- Borrowing creates a temporary alias
- The borrowed-from place is **unusable** as long as the borrow exists



- The Rust compiler uses a static analysis to determine when borrows expire
- Borrow checking across functions may require lifetime annotations (not discussed here)

Viper encoding: references

```
struct Pair {  
  first: i32,  
  second: i32,  
}
```



```
field ref: Ref  
predicate RefMutPair(this: Ref) {  
  acc(this.ref) && Pair(this.ref)  
}
```

```
let mut p;  
p = Pair{ first: 5, second: 7 };  
let q = &mut p;  
q.first = 9;  
// borrow for q expires  
swap(&mut p);
```



```
...  
q := new(ref)  
q.ref := p  
fold RefMutPair(q)  
...  
unfold RefMutPair(q)  
unfold Pair(q.ref)  
q.ref.first := makeInt(9)  
fold Pair(q.ref)  
fold RefMutPair(q)  
...  
unfold RefMutPair(q)  
...
```

Functional correctness


- The [automatically-generated core proof](#) re-proves memory safety
- The contained permission specifications enable reasoning about [aliasing](#) and [framing](#) (and race freedom in concurrent code)
- Functional specifications can be layered on top easily

```
#[ensures(p.first == old(p.second))]  
#[ensures(p.second == old(p.first))]  
fn swap(p: &mut Pair) {  
    let tmp = p.first;  
    p.first = p.second;  
    p.second = tmp;  
}
```

Reborrowing

- Reborrowing occurs when a function borrows from a borrowed value

```
fn getFirst(p: &mut Pair) -> &mut i32 {  
    &mut p.first  
}
```



```
method getFirst(p: Ref) returns (res: Ref)  
    requires RefMutPair(p)  
    ensures RefMutI32(res)  
{ ... }
```

```
fn demo(mut p: Pair) {  
    p.second = 0;  
    let f = getFirst(&mut p);  
    *f = *f + 1;  
    // borrow for f expires  
    assert!(p.second == 0);  
}
```


How to re-construct
RefMutPair(p)?

Naïve encoding does not
return permissions to the
remainder of the Pair p

Viper encoding: reborrowing

- Reborrowing creates a partial data structure

```
fn getFirst(p: &mut Pair) -> &mut i32 {  
  &mut p.first  
}
```

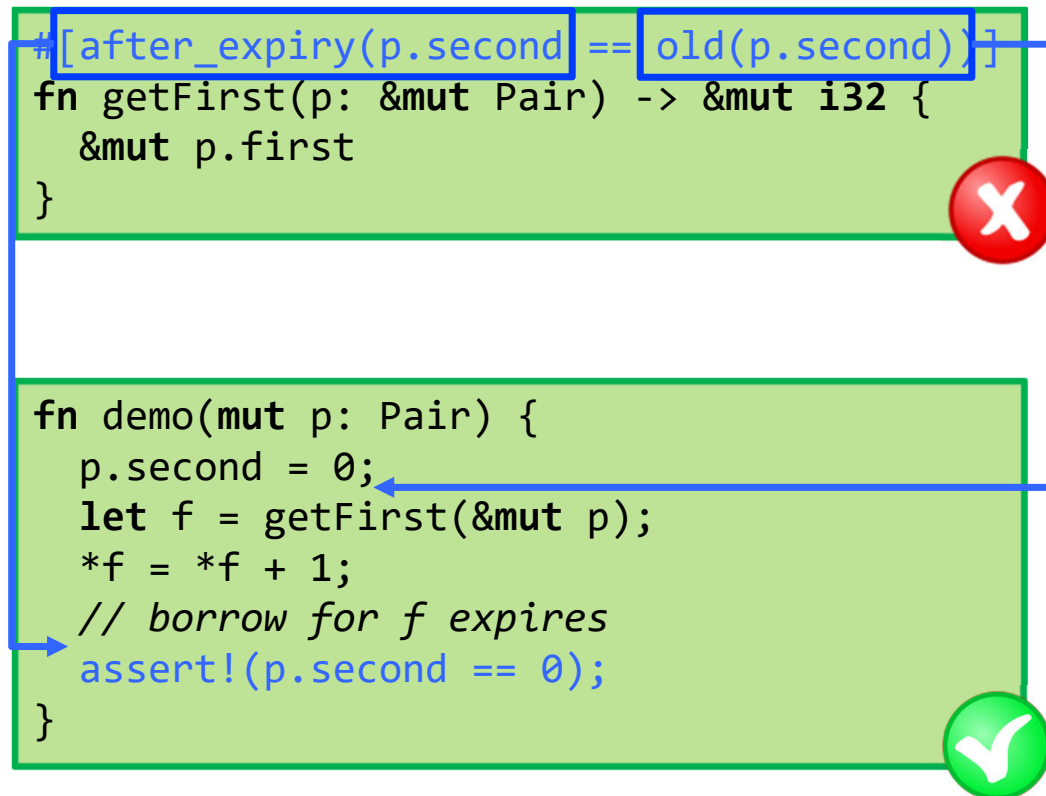


- Partial data structures can be described using magic wands

```
method getFirst(p: Ref) returns (res: Ref)  
  requires RefMutPair(p)  
  ensures RefMutI32(res)  
  ensures RefMutI32(res) --* RefMutPair(p)  
  { ... }
```

- Creating a reborrow packages a magic wand, expiring a reborrow applies it

Functional correctness: pledges




- Lender function could modify its argument before creating the reborrow
- Postconditions **must not access unusable borrowed-from value**
 - All assertions follow Rust's type rules
- **Pledges** specify properties that **will be true when the reborrow expires**
- Here, `getFirst` can predict the value of `p.second` at expiry because the caller cannot modify this unusable place before the reborrow expires

Fun with pledges

```
#[ensures(*result == old(p.first))]  
#[after_expiry(p.first ==  
               before_expiry(*result))]  
fn getFirst(p: &mut Pair) -> &mut i32 {  
    &mut p.first  
}
```

```
fn demo(mut p: Pair) {  
    p.first = 0;  
    let f = getFirst(&mut p);  
    *f = *f + 1;  
    // borrow for f expires  
    assert!(p.first == 1);  
}
```



- This pledge does not allow the client to modify the (usable) reborrowed value
- The postcondition **can specify** side effects on the reborrowed result
- **Pledges** specify properties of the borrowed-from value parametric in the reborrowed value

Viper encoding: pledges

- Pledges are conjoined to the rhs of the magic wand for the reborrow
 - Analogous to other functional specifications

```
#[after_expiry(p.second == old(p.second))]  
fn getFirst(p: &mut Pair) -> &mut i32 {  
    &mut p.first  
}
```



```
method getFirst(p: Ref) returns (res: Ref)  
    requires RefMutPair(p)  
    ensures RefMutI32(res)  
    ensures RefMutI32(res) --* RefMutPair(p) &&  
        p.ref.second == old(p.ref.second)  
{ ... }
```

- When the reborrow expires, applying the wand provides the properties from the pledge

Shared references

- Rust distinguishes between exclusive, mutable references and shared, immutable references

- Allows programmers to express design intent
- Enables race-free concurrent reads

```
fn getFirst(p: &Pair) -> &i32 {  
    &p.first  
}
```

- Shared borrowing creates a temporary alias
- The borrowed-from place remains usable, but immutable as long as a shared borrow exists

Shared borrow
Creating another
shared reference
p remains usable
but Type is
q1 is immutable
Borrow expires
again

```
let mut p = Pair{ first: 5, second: 7 };  
let q1 = &p;  
let q2 = &p;  
let v = p.first;  
q1.first = 9;  
swap(&mut p);
```


Aliasing XOR mutability

```
fn getFirst(p: &Pair) -> &i32 {  
    &p.first  
}
```

```
fn demo(mut p: Pair) {  
    p.first = 0;  
    let f = getFirst(&p);  
    // borrow for f expires  
    assert!(p.first == 0);  
}
```



- Rust guarantees that, in each state, there is either **at most one usable place** that can access a value, **or the value is immutable**
- This guarantee allows us **to frame values automatically**
- No value can change while a shared reference exists

Viper encoding: shared references

- We reuse the type encoding for mutable references


```
struct Pair {  
  first: i32,  
  second: i32,  
}
```



```
field ref: Ref  
predicate RefMutPair(this: Ref) {  
  acc(this.ref) && Pair(this.ref)  
}
```

- and use **fractional permissions** to express shared access
 - read is a symbolic positive permission amount

```
fn isSame(p: &Pair) -> bool {  
  p.first == p.second  
}
```



```
method isSame(p: Ref) returns (res: Bool)  
  requires acc(RefMutPair(p), read)  
  ensures pBool(res)  
  { ... }
```

Viper encoding: creating and expiring shared references

```
let mut p;  
p = Pair{ first: 5, second: 7 };  
let q = &p;  
let v = q.first;  
// borrow for q expires  
swap(&mut p);
```



```
...  
q := new()  
exhale acc(Pair(p), write-read)  
inhale acc(RefMutPair(q), read)  
assume unfolding acc(RefMutPair(q), read) in  
  q.ref == p  
...  
v := unfolding acc(RefMutPair(q), read) in  
  unfolding acc(Pair(q.ref), read) in  
    q.ref.first  
...  
exhale acc(RefMutPair(q), read)  
inhale acc(Pair(p), write-read)  
...
```

- This encoding *trusts the borrow checker* to determine when all shared borrows have expired, which simplifies permission accounting

Comparison of annotation overhead

```
#[feature(box_patterns)]

use prusti_contracts::*;

struct Node {
    elem: i32,
    next: List,
}

enum List {
    Empty,
    More(Box<Node>),
}

impl List {
    #[pure]
    #[ensures(result >= 0)]
    fn len(&self) -> usize {
        match self {
            List::Empty => 0,
            List::More(box node) =>
                1 + node.next.len(),
        }
    }

    #[ensures(result.len() ==
        self.len() + that.len())]
    pub fn zip(&self, that: &List) -> List {
        match self {
            List::Empty => that.cloneList(),
            List::More(box node) => {
                let new_node = Box::new(Node {
                    elem: node.elem,
                    next: that.zip(&node.next),
                });
                List::More(new_node)
            }
        }
    }
}
```

```
#[ensures(result.len() == self.len())]
pub fn cloneList(& self) -> List {
    match self {
        List::Empty => List::Empty,
        List::More(box node) => {
            let new_node = Box::new(Node {
                elem: node.elem,
                next: node.next.cloneList(),
            });
            List::More(new_node)
        }
    }
}
```

P*rust-*i

```
field next: Ref
field elem: Int

predicate list(this: Ref) {
    acc(this.elem) && acc(this.next) &&
    (this.next != null => list(this.next))
}

function len(this: Ref): Int
    requires acc(list(this), wildcard)
{
    unfolding acc(list(this), wildcard) in
    (this.next == null ? 0 : len(this.next) + 1)
}

method zip(this: Ref, that: Ref)
    returns (res: Ref)
    requires acc(list(this), 1/2) &&
        acc(list(that), 1/2)
    ensures acc(list(this), 1/2) &&
        acc(list(that), 1/2)
    ensures list(res)
    ensures res != null
    ensures len(res) == len(this) + len(that)
{
    unfold acc(list(this), 1/2)
    if(this.next == null) {
        res := cloneList(that)
    } else {
        res := new(*)
        res.elem := this.elem
        var rest: Ref
        rest := zip(that, this.next)
        res.next := rest
        fold list(res)
    }
    fold acc(list(this), 1/2)
}
```


```
method cloneList(this: Ref) returns (res: Ref)
    requires acc(list(this), 1/2)
    ensures acc(list(this), 1/2) && list(res)
    ensures res != null
    ensures len(res) == len(this)
{
    res := new(*)
    unfold acc(list(this), 1/2)
    if(this.next == null) {
        res.next := null
    } else {
        var tmp: Ref
        tmp := cloneList(this.next)
        res.elem := this.elem
        res.next := tmp
    }
    fold acc(list(this), 1/2)
    fold list(res)
}
```

VIPER

A heap-free encoding: purification

- The behavior of a Rust function can be described entirely via its parameter and result **values**
- This suggests an alternative encoding (purification) that does not use the heap

```
struct Pair {  
  first: i32,  
  second: i32,  
}
```



```
adt Pair {  
  MkPair(first: Int, second: Int)  
}
```

```
#[ensures(p.first == old(p.second))]  
#[ensures(p.second == old(p.first))]  
fn swap(p: &mut Pair) {  
  let tmp = p.first;  
  p.first = p.second;  
  p.second = tmp;  
}
```



```
method swap(p: Pair) returns (res: Pair)  
  ensures res.first == p.second  
  ensures res.second == p.first  
{  
  res := p  
  var tmp: Int := res.first  
  res := MkPair(res.second, res.second)  
  res := MkPair(res.first, tmp)  
}
```

Purification

- Purification simplifies the verification problem substantially
 - No heap and, thus, no need for permissions, predicates, magic wands
 - Much faster verification times (Creusot and Verus verifiers)
 - But treatment of reborrows is not trivial (based on prophecy variables)
- However, Rust allows **unsafe code** to work around the limitations of the strict ownership type system
 - Unsafe code may introduce aliasing among mutable pointers

```
let mut p;  
p = Pair{ first: 5, second: 7 };  
let ptr1: *mut i32 = &mut p.first;  
let ptr2 = ptr1;  
unsafe { *ptr1 = *ptr1 + 1; }  
unsafe { *ptr2 = *ptr2 + 1; }  
println!("{}", p.first);
```

- ptr1 and ptr2 are mutable aliases
- Verification requires a logic that can handle general heaps

Summary: Rust verification in Prusti

```
#[ensures(*x == old(*y) )]  
#[ensures(*y == old(*x) )]  
fn swap(x: &mut i32, y: &mut i32) {  
    let tmp = *x;  
    *x = *y;  
    *y = tmp;  
}
```

- Prusti extracts permissions (and predicates) automatically from type information
- A Viper “core proof” of memory safety is generated completely automatically
- Users can add **functional correctness** specifications, by using a slight extension of Rust expressions

The overhead for programmers is substantially reduced
(both amount and complexity of annotations)

Outline

- The Viper intermediate verification language
 - Reasoning about the heap
 - Abstraction
 - Advanced separation logic in Viper
- Verification of Rust programs
 - Ownership
 - References
- Conclusion

Challenges revisited

Heap data structures pose three major challenges for sequential verification

- Reasoning about aliasing
 - Permissions and separating conjunction
- Framing, especially for dynamic data structures
 - Sound frame rule, predicates
- Writing specifications that preserve information hiding
 - Data abstraction, heap-dependent functions



Ownership in program verification

Ownership types
[Clarke, Potter, Noble '98]

Universe types
[Müller et al. '99]

Ownership-based verification
[Müller '01]

Separation Logic
[O'Hearn, Reynolds, Yang '01]

Spec# [Leino, Müller '04]

Rust [Matsakis, Klock '14]

Prusti [Astrauskas et al. '19]
Creusot [Denis et al. '22]
Verus [Latuada et al. '23]

Move [Blackshear et al. '19]

Move Prover [Zhong et al. '20]

